

Abhishek Mishra

Solutions Architect — AI Solutions & Intelligent Data Platforms

Email mishraabhishek.mishra021@gmail.com LinkedIn [linkedin.com/in/abhi021](https://www.linkedin.com/in/abhi021) GitHub github.com/abhishek021 Location Bengaluru, India

SUMMARY PROFILE

Solutions Architect with 8 years of experience evolving from hands-on data engineering into platform and solution architecture. I design scalable, governed data and AI platforms spanning the full lifecycle — data engineering, AI/ML enablement, and the operational foundations (DataOps, MLOps, LLMOps, AIOps) that keep enterprise platforms reliable in production.

Today I architect enterprise-grade AI solutions and intelligent data platforms on Databricks and Spark for global CPG and pharma clients, while bringing agentic systems into production to drive self-healing, zero-touch operations. My core strength is converting complex data ecosystems into governed, observable, production-ready platforms — and enabling the engineers I work with to grow alongside them.

EXPERIENCE SIGMOID · 8 YRS

- Solutions Architect

Sigmoid · Bengaluru, India

Architecting enterprise-grade AI solutions and intelligent data platforms, and reimagining platform engineering with agentic systems. Own the design of scalable, resilient platforms across data engineering, AI engineering, and DataOps — embedding governance, observability, and automation, and leveraging agentic AI to drive self-healing, zero-touch operations.

Jan 2026 – Present
- Data Operations & Engineering Lead

Sigmoid · Bengaluru, India

- Led a team of 18+ engineers, owning delivery across the full data lifecycle — from ingestion and pipeline reliability to production support and platform governance.
 - Drove a large-scale platform modernization on Databricks (Unity Catalog, Delta Live Tables, Asset Bundles), establishing governed, observable, production-grade standards.
 - Built the team from the ground up — defining hiring bars, onboarding engineers, and shaping the technical practices that scaled delivery.
 - Introduced GenAI and LLMOps into operational workflows, applying LLMs and agentic patterns to improve automation, triage, and self-healing.
 - Partnered with L&D to design and run a structured intern-to-full-time program — mentoring and assessing 40+ interns from hiring and technical assessment through training to conversion.

Jan 2024 – Dec 2025
- Senior Data Operations Engineer

Sigmoid · Bengaluru, India

- Designed and operated production pipelines on Databricks, applying data modeling and engineering best practices to improve reliability and maintainability.
 - Implemented security architecture and access-governance controls across data platforms, strengthening compliance and auditability.
 - Drove automation across operational workflows, reducing manual intervention and improving pipeline stability and recovery.
 - Mentored junior engineers, establishing standards that raised the quality bar for pipeline development.

Jan 2023 – Dec 2023
- DataOps Engineer II

Sigmoid · Bengaluru, India

- Migrated Python pipelines to Databricks, leveraging ACID transactions, schema evolution, merge/update/delete DML, scalable metadata, and delta-log auditing.
 - Optimized PySpark and Delta workloads over billions of rows using partitioning, OPTIMIZE, and time travel.
 - Orchestrated Databricks jobs from Airflow using job and all-purpose clusters; managed AWS web servers with performance monitoring.

Aug 2021 – Dec 2022
- DataOps Engineer I

Sigmoid · Bengaluru, India

- Set up and managed AWS Hadoop clusters running TB-scale ETL pipelines for multiple customers.
 - Implemented zero-trust security, bastion-host architecture, and Grafana-based log auditing with continuous monitoring of AWS resources.
 - Ran PoCs that brought new technology into the stack (Alluxio caching, Databricks Delta Lake); deployed Hadoop, Kubernetes monitoring, and GitLab estates.

Mar 2019 – Sep 2021
- Associate DataOps Engineer

Sigmoid · Bengaluru, India

- Deployed and managed GCP and AWS Hadoop clusters running Spark ETL jobs.
 - Wrote scripts to manage and monitor thousand-node clusters running 24/7 real-time pipelines; built automation for productivity and cost reduction.

Jul 2018 – Feb 2019

KEY PROJECTS AGENTIC AI

- Multi-Agent Intent System

LangGraph agents handling distinct user intents with runtime routing, session state, conversation history in a relational DB, multi-step workflows via graphs/subgraphs, full LangSmith traceability, a monitoring dashboard, and CI/CD with automated LLM evaluation.
- Permission-Aware Knowledge Assistant

Natural-language access across pharma SOPs, dossiers, and quality records with identity-scoped retrieval, so each user sees only what their role permits.

Natural-Language Sales Analytics – Databricks Genie semantic layer (ontology, measures, business logic) letting CPG brand and category managers query in plain English with one trusted definition of metrics.

Self-Healing Data Pipeline — Deterministic supervisor plus LLM repair agents with policy-driven schema governance and Bronze immutability; pipelines classify drift, repair safely, and halt for review on policy violations.

SKILLS

DATA & AI PLATFORM

Databricks

Unity Catalog

Delta Live Tables

Asset Bundles

Agent Bricks

Databricks Genie

Spark

PySpark

Airflow

AGENTIC AI & ML

LangChain

LangGraph

LangSmith

RAG

MCP

MLflow

GenAI

MLOps

LLMOps

CLOUD & DEVOPS

AWS

Azure

GCP

CI/CD

Azure DevOps

Kubernetes

Terraform

Docker

Grafana

LANGUAGES & FOUNDATIONS

Python

SQL

Shell

Systems Design

Enterprise Architecture

Security Architecture

LEADERSHIP

Lead a team of **18+ engineers**. In partnership with L&D, have driven the intern-to-full-time journey end to end — mentoring, assessing, and evaluating **40+ interns** from hiring and technical assessment through training to full-time conversion.

EDUCATION

2014 – 2018

B.E. Computer Engineering — 77%
University Institute of Technology, The University of Burdwan

LANGUAGES

English (Full Professional) · Hindi (Native / Bilingual) · Bengali